

## 1.2 Approaching Coding Questions

Coding interviews typically last 30 to 45 minutes and come in a variety of formats. Early in the interview process, coding interviews are often conducted via remote coding assessment tools like HackerRank, Codility, or CoderPad. During final-round onsite interviews, it's typical to write code on a whiteboard. Regardless of the format, the approach outlined below to solve coding interview problems applies.

1

**After receiving the problem:** Don't jump right into coding. It's crucial first to make sure you are solving the correct problem. Due to language barriers, misplaced assumptions, and subtle nuances that are easy to miss, misunderstanding the problem is a frequent occurrence. To prevent this, make sure to repeat the question back to the interviewer so that the two of you are on the same page. Clarify any assumptions made, like the input format and range, and be sure to ask if the input can be assumed to be non-null or well formed. As a final test to see if you've understood the problem, work through an example input and see if you get the expected output. Only after you've done these steps are you ready to begin solving the problem.

2

**When brainstorming a solution:** First, explain at a high level how you could tackle the question. This usually means discussing the brute-force solution. Then, try to gain an intuition for why this brute-force solution might be inefficient, and how you could improve upon it. If you're able to land on a more optimal approach, articulate how and why this new solution is better than the first brute-force solution provided. Only after you've settled on a solution is it time to begin coding.

3

**When coding the solution:** Explain what you are coding. Don't just sit there typing away, leaving your interviewer in the dark. Because coding interviews often let you pick the language you write code in, you're expected to be proficient in the programming language you chose. As such, avoid pseudocode in favor of proper compilable code. While there is time pressure, don't take many shortcuts when coding. Use clear variable names and follow good code organization principles. Write well-styled code — for example, following PEP 8 guidelines when coding in Python. While you are allowed to cut some corners, like assuming a helper method exists, be explicit about it and offer to fix this later on.

4

**After you're done coding:** Make sure there are no mistakes or edge cases you didn't handle. Then write and execute test cases to prove you solved the problem.

At this point, the interviewer should dictate which direction the interview heads. They may ask about the time and space complexity of your code. Sometimes they may ask you to refactor and clean the code, especially if you cut some corners while coding the solution. They may also extend the problem, often with a new constraint. For example, they may ask you not to use recursion and instead tell you to solve the problem recursively. Or, they might ask you to not use surplus memory and instead solve the problem in place. Sometimes, they may pose a tougher variant of the problem as a follow-up, which might require starting the problem-solving process all over again.

Start

Understand the  
problem w/ a  
toy example

Brainstorm the  
brute-force version  
of the soln.

Proper code ethics  
while programming  
like PEP8 ethics

- Solve for all types of edge cases.
- Discuss time & space complexity
- Refactor/clean the code

X

Solve against all types  
of constraints:

- Don't use recursion but solve iteratively.
- Don't use variables necessarily and solve it in-place.
- Tougher variants of the problem could make some be revisited.

Stop

## **Replacement**

Suppose in a given text  $T$  we want to replace the first occurrence of a pattern  $P_1$  by a pattern  $P_2$ . We will denote this operation by

REPLACE(text, pattern<sub>1</sub>, pattern<sub>2</sub>)

For example

C C      1st occurrence gets replaced only.

REPLACE(XABYABZ, 'AB', 'C') = 'XCYABZ'  
REPLACE(XABYABZ, 'BA', 'C') = 'XABYABZ'

In the second case, the pattern BA does not occur, and hence there is no change.

We note that this REPLACE function can be expressed as a deletion followed by an insertion if we use the preceding DELETE and INSERT functions. Specifically, the REPLACE function can be executed by using the following three steps:

K := INDEX(T, P<sub>1</sub>)  
T := DELETE(T, K, LENGTH(P<sub>1</sub>))  
INSERT(T, K, P<sub>2</sub>)

The first two steps delete P<sub>1</sub> from T, and the third step inserts P<sub>2</sub> in the position K from which P<sub>1</sub> was deleted.

Suppose a text T and patterns P and Q are in the memory of a computer. Suppose we want to replace every occurrence of the pattern P in T by the pattern Q. This might be accomplished by repeatedly applying

REPLACE(T, P, Q)

until INDEX(T, P) = 0 (i.e., until P does not appear in T). An algorithm which does this follows.

## 3.5 STRING OPERATIONS

Although a string may be viewed simply as a sequence or linear array of characters, there is a fundamental difference in use between strings and other types of arrays. Specifically, groups of consecutive elements in a string (such as words, phrases and sentences), called substrings, may be units unto themselves. Furthermore, the basic units of access in a string are usually these substrings, not individual characters.

Consider, for example, the string

TO BE OR NOT TO BE'

We may view the string as the 18-character sequence T, O, □, B, ..., E. However, the substrings TO, BE, OR ... have their own meaning.

On the other hand, consider an 18-element linear array of 18 integers,

4, 8, 6, 15, 9, 5, 4, 13, 8, 5, 11, 9, 9, 13, 7, 10, 6, 11

The basic unit of access in such an array is usually an individual element. Groups of consecutive elements normally do not have any special meaning.

For the above reason, various string operations have been developed which are not normally used with other kinds of arrays. This section discusses these string-oriented operations. The next section shows how these operations are used in word processing. (Unless otherwise stated or implied, we assume our character-type variables are dynamic and have a variable length determined by the context in which the variable is used.)

## ① Substring ② position of substring's last character

Accessing a substring from a given string requires three pieces of information: (1) the name of the string or the string itself, (2) the position of the first character of the substring in the given string and (3) the length of the substring or the position of the last character of the substring. We call this operation SUBSTRING. Specifically, we write

`SUBSTRING(string, initial, length)`

to denote the substring of a string  $S$  beginning at position  $K$  and having a length  $L$ .

## ③ Length of substring

## ② Indexing

- Indexing, also called *pattern matching*, refers to finding the position where a string pattern P first appears in a given string text. We call this operation INDEX and write  
 $\text{INDEX}(\text{text}, \text{pattern})$
- If the pattern P does not appear in the text L, then INDEX is assigned the value 0. The arguments "text" and "pattern" can be either string constants or string variables.

Observe the reverse order of the arguments in UCSD Pascal:

## ③ Concatenation

Let  $S_1$  and  $S_2$  be strings. Recall (Sec. 3.2) that the concatenation of  $S_1$  and  $S_2$ , which we denote by  $S_1 \# S_2$ , is the string consisting of the characters of  $S_1$  followed by the characters of  $S_2$ .

## ④ Length

Context of a character  
not of entire space

The number of characters in a string is called its length. We will write

for the length of a given string. Thus  
 $\text{LENGTH}(\text{COMPUTER}) = 8$

$\text{LENGTH}(\text{string})$   
 $\text{LENGTH}[0] = 0$   
 $\text{LENGTH}[1] = 1$  white space  
 $\text{LENGTH}[2] = 2$  character  
 $\text{LENGTH}[3] = 3$

Some of the programming languages denote this function as follow:

PL/I:  $\text{LENGTH(string)}$   
BASIC:  $\text{LEN(string)}$   
UCSD Pascal:  $\text{LENGTH(string)}$   
SNOBOL:  $\text{SIZE(string)}$

FORTRAN and standard Pascal, which use fixed-length string variables, do not have any built-in LENGTH functions for strings. However, such variables may be viewed as having variable length if one ignores all trailing blanks. Accordingly, one could write a subprogram LENGTH in these languages so that

Historically, computers were first used for processing numerical data. Today, computers are frequently used for processing nonnumerical data, called character data. This chapter discusses how such data are stored and processed by the computer.

One of the primary applications of computers today is in the field of word processing. Such processing usually involves some type of pattern matching, as in checking to see if a particular word S appears in a given text T. We discuss this pattern matching problem in detail and, moreover, present two different pattern matching algorithms. The complexity of these algorithms is also investigated. Computer terminology usually uses the term "string" for a sequence of characters rather than the term "word," since "word" has another meaning in computer science. For this reason, many texts sometimes use the expression "string processing," "string manipulation" or "text editing" instead of the expression "word processing."

6. In C++, an array index starts with 0.
7. An array index can be any expression that evaluates to a nonnegative integer. The value of the index must always be less than the size of the array.
8. There are no aggregate operations on arrays, except for the input/output of character arrays (C-strings).
9. Arrays can be initialized during their declaration. If there are fewer initial values than the array size, the remaining elements are initialized to 0.
10. The base address of an array is the address of the first array component. For example, if `list` is a one-dimensional array, the base address of `list` is the address of `list[0]`.

11. When declaring a one-dimensional array as a formal parameter, you usually omit the array size. If you specify the size of a one-dimensional array in the formal parameter declaration, the compiler will ignore the size.
  12. In a function call statement, when passing an array as an actual parameter, you use only its name.
13. As parameters to functions, arrays are passed by reference only.
14. Because as parameters, arrays are passed by reference only, when declaring an array as a formal parameter, you do not use the symbol & after the data type.
15. A function cannot return a value of type array.

## 13.5.4 Search Engine Indexing

The World Wide Web contains a huge collection of text documents (Web pages). Information about these pages are gathered by a program called a **Web crawler**, which then stores this information in a special dictionary database. A **Web search engine** allows users to retrieve relevant information from this database, thereby identifying relevant pages on the Web containing given keywords. In this section, we present a simplified model of a search engine.

## Inverted Files

The core information stored by a search engine is a dictionary, called an ***inverted index*** or ***inverted file***, storing key-value pairs  $(w, L)$ , where  $w$  is a word and  $L$  is a collection of pages containing word  $w$ . The keys (words) in this dictionary are called ***index terms*** and should be a set of vocabulary entries and proper nouns as large as possible. The elements in this dictionary are called ***occurrence lists*** and should cover as many Web pages as possible.

We can efficiently implement an inverted index with a data structure consisting of the following:

1. An array storing the occurrence lists of the terms (in no particular order);
2. A compressed trie for the set of index terms, where each leaf stores the index of the occurrence list of the associated term.

For more details, see [this article](#).

A **trie** (pronounced “try”) is a tree-based data structure for storing strings in order to support fast pattern matching. The main application for tries is in information retrieval. Indeed, the name “trie” comes from the word “retrieval.” In an information retrieval application, such as a search for a certain DNA sequence in a genomic database, we are given a collection  $S$  of strings, all defined using the same alphabet. The primary query operations that tries support are pattern matching and **prefix matching**. The latter operation involves being given a string  $X$ , and looking for all the strings in  $S$  that contain  $X$  as a prefix.

Data Structure	Time Complexity						Space Complexity
	Average	Search	Insertion	Deletion	Access	Worst	
Array	$O(1)$	$O(n)$	$O(n)$	$O(n)$	$O(1)$	$O(n)$	$O(n)$
Stack	$O(n)$	$O(n)$	$O(1)$	$O(1)$	$O(n)$	$O(n)$	$O(n)$
Queue	$O(n)$	$O(n)$	$O(1)$	$O(1)$	$O(n)$	$O(n)$	$O(n)$
Linked List	$O(n)$	$O(n)$	$O(1)$	$O(1)$	$O(n)$	$O(n)$	$O(n)$
Hash Map	N/A	$O(1)$	$O(1)$	N/A	$O(n)$	$O(n)$	$O(n)$
Binary Search Tree	$O(\log(n))$	$O(\log(n))$	$O(\log(n))$	$O(\log(n))$	$O(n)$	$O(n)$	$O(n)$

## Arrays

An array is a series of consecutive elements stored sequentially in memory. Arrays are optimal for accessing elements at particular indices, with an  $O(1)$  access and index time. However, they are slower for searching and deleting a specific value, with an  $O(N)$  runtime, unless sorted. An array's simplicity makes it one of the most commonly used data structures during coding interviews.

HashMap: The HashMap collection is widely used, both in interviews and in the real world. We've provided a snippet of the syntax below.

```
1 HashMap<String, String> map = new HashMap<String, String>();  
2 map.put("one", "uno");  
3 map.put("two", "dos");  
4 System.out.println(map.get("one"));
```

Before your interview, make sure you're very comfortable with the above syntax. You'll need it.

Java's collection framework is incredibly useful, and you will see it used throughout this book. Here are some of the most useful items:

ArrayList: An ArrayList is a dynamically resizing array, which grows as you insert elements.

```
1 ArrayList<String> myArr = new ArrayList<String>();  
2 myArr.add("one");  
3 myArr.add("two");  
4 System.out.println(myArr.get(0)); /* prints <one> */
```

Vector: A vector is very similar to an ArrayList, except that it is synchronized. Its syntax is almost identical as well.

```
1 Vector<String> myVect = new Vector<String>();  
2 myVect.add("one");  
3 myVect.add("two");  
4 System.out.println(myVect.get(0));
```

LinkedList: LinkedList is, of course, Java's built-in `LinkedList` class. Though it rarely comes up in

### 3.1 Contiguous vs. Linked Data Structures

Data structures can be neatly classified as either contiguous or linked, depending upon whether they are based on arrays or pointers:

- Contiguously-allocated structures are composed of single slabs of memory, and include arrays, matrices, heaps, and hash tables.
- Linked data structures are composed of distinct chunks of memory bound together by pointers, and include lists, trees, and graph adjacency lists.

together by pointers, and include lists, trees, and graph adjacency lists.

In this section, we review the relative advantages of contiguous and linked data structures. These tradeoffs are more subtle than they appear at first glance, so I encourage readers to stick with me here even if you may be familiar with both types of structures.

### 3.1.1 Arrays

### 3.1.1 Arrays

The array is the fundamental contiguously-allocated data structure. [Arrays are structures of fixed-size data records such that each element can be efficiently located by its index or (equivalently) address.]

A good analogy likens an array to a street full of houses, where each array element is equivalent to a house, and the index is equivalent to the house number. Assuming all the houses are equal size and numbered sequentially from 1 to  $n$ , we can compute the exact position of each house immediately from its address.<sup>1</sup>

Advantages of contiguously-allocated arrays include:

- Constant-time access given the index - Because the index of each element maps directly to a particular memory address, we can access arbitrary data items instantly provided we know the index.
- Space efficiency - Arrays consist purely of data, so no space is wasted with links or other formatting information. Further, end-of-record information is not needed because arrays are built from fixed-size records.
- Memory locality - A common programming idiom involves iterating through all the elements of a data structure. Arrays are good for this because they exhibit excellent memory locality. Physical continuity between successive data accesses helps exploit the high-speed *cache memory* on modern computer architectures.

{ The downside of arrays is that we cannot adjust their size in the middle of a program's execution. Our program will fail soon as we try to add the  $(n +$

Houses in Japanese cities are traditionally numbered in the order they were built, not by their physical location. This makes it extremely difficult to locate a Japanese address without a detailed map.

1)st customer, if we only allocate room for  $n$  records. We can compensate by allocating extremely large arrays, but this can waste space, again restricting what our programs can do.]

Actually, we can efficiently enlarge arrays as we need them, through the miracle of dynamic arrays. Suppose we start with an array of size 1, and double its size from  $m$  to  $2m$  each time we run out of space. This doubling process involves allocating a new contiguous array of size  $2m$ , copying the contents of the old array to the lower half of the new one, and returning the space used by the old array to the storage allocation system.

The apparent waste in this procedure involves the recopying of the old contents on each expansion. How many times might an element have to be recopied after a total of  $n$  insertions? Well, the first inserted element will have been recopied when the array expands after the first, second, fourth, eighth, ... insertions. It will take  $\log_2 n$  doublings until the array gets to have  $n$  positions. However, most elements will do not suffer much upheaval. Indeed, the  $(n/2 + 1)$ st through  $n$ th elements will move at most once and might never have to move at all.

If half the elements move once, a quarter of the elements twice, and so on, the total number of movements  $M$  is given by

$$M = \sum_{i=1}^{\lg n} i \cdot n / 2^i = n \sum_{i=1}^{\infty} i / 2^i \leq n \sum_{i=1}^{\infty} i / 2^i = 2n$$

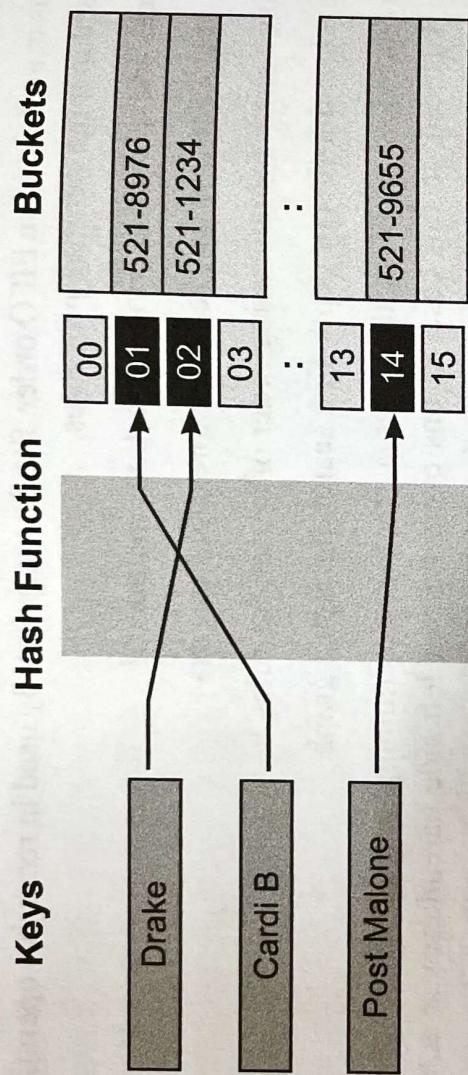
*partial sum*      *larger as sum*

*infinite series*      *converges to 2*

Thus, each of the  $n$  elements move only two times on average, and the total work of managing the dynamic array is the same  $O(n)$  as it would have been if a single array of sufficient size had been allocated in advance! The primary thing lost using dynamic arrays is the guarantee that each array access takes constant time in the worst case. Now all the queries will be fast, except for those relatively few queries triggering array doubling. What we get instead is a promise that the  $n$ th array access will be completed quickly enough that the total effort expended so far will still be  $O(n)$ . Such amortized guarantees arise frequently in the analysis of data structures.

## Hash Maps

A hash map stores key-value pairs. For every key, a hash map uses a hash function to compute an index, which locates the bucket where that key's corresponding value is stored. In Python, a dictionary offers support for key-value pairs and has the same functionality as a hash map.



While a hash function aims to map each key to a unique index, there will sometimes be “collisions” where different keys have the same index. In general, when you use a good hash function, expect the elements to be distributed evenly throughout the hash map. Hence, lookups, insertions, or deletions for a key take constant time.

Due to their optimal runtime properties, hash maps make a frequent appearance in coding interview questions.